

L9: Deep Neural Network

Shan Wang

Lingnan College, Sun Yat-sen University

2020 Data Mining and Machine Learning LN3119

<https://wangshan731.github.io/DM-ML/>



Last lecture

- One-layer Perceptron

- $o(\mathbf{x}) = f(x_1w_1 + x_2w_2 + x_3w_3 + b)$

- Multi-layer Perceptron

- Model: input, hidden, output

- Strategy: minimize error

- Algorithm: BP algorithm

- $w_{i,j}' = w_{i,j} - \eta \varepsilon_i h_j$


- $\varepsilon_i = \sum_{k=1}^K \varepsilon_k w_{k,i} f'^i$

- Overfitting

- Application

Forward Propagation of Info.

Backward Propagation of Error



Course Outline

- Supervised learning
 - Linear regression
 - Logistic regression
 - SVM and kernel
 - Tree models
- Deep learning
 - Neural networks
 - Convolutional NN
 - Recurrent NN
- Unsupervised learning
 - Clustering
 - PCA
 - EM
- Reinforcement learning
 - MDP
 - ADP
 - Deep Q-Network

This lecture

- Deep Learning
- Deep Auto Encoder
- Convolutional NN
- Recurrent NN

Deep Learning

Deep Learning

- “Deep”: the structure is deep, contains many layers
- “Learning”
 - Supervised learning: input data has label
 - Multi-Layer Perceptron, CNN, RNN
 - Unsupervised learning: input data has no label
 - Deep Auto-Encoder
 - Reinforce learning: use penalty and reward
 - FNN and RNN applied in RL, Deep Q-Network

Why we need “deep”?


- With more hidden layers, NNs are expected to describe the reality better
 - Functions that can be compactly represented by a depth K architecture, might require an exponential number of neurons to be represented by a depth $K-1$ architecture
 - Successive layers can learn higher-level features

Problems from Depth

- Lack of big **data**
 - Now we have a lot of big data
- Lack of **computational** resources
 - Now we have GPUs and HPCs
- **Local optimality**
 - Add momentum, pre-training techniques & various optimization algorithms
- **Gradient vanishing**
 - Use ReLU activation function...
- Too many **parameters**
 - Train the network layer by layer & dropout



Multi-Layer Perceptron

- MLP is a **fully connected** neural network with multiple hidden layers

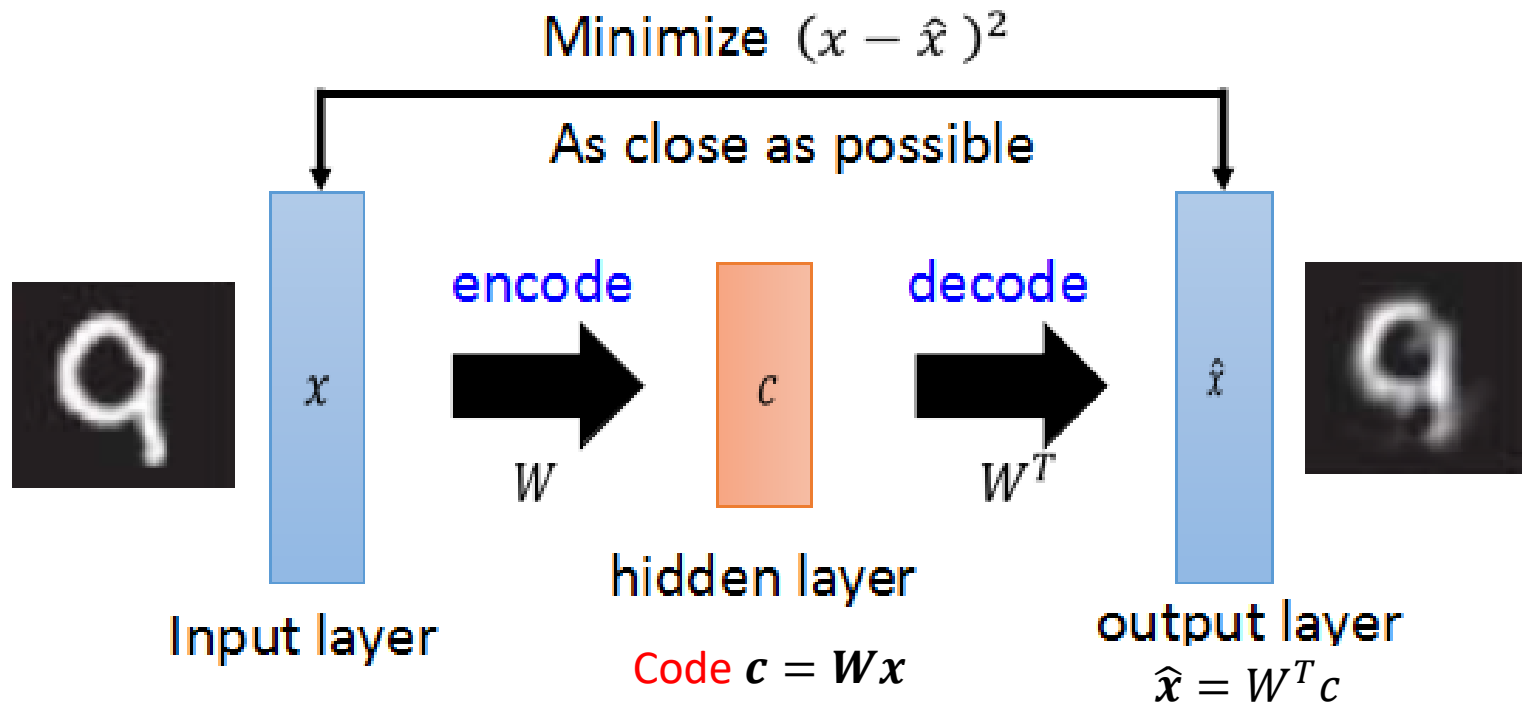


Deep Auto-Encoder

Auto-Encoder


- Auto-Encoder can be used for data suppress, dimensionality reduction, pre-training NNs etc.
- An auto-encoder contains an **encoder** and a **decoder**
- The encoder plays the role of coding the original data
 - Data A  Codes
- The decoder plays the role of decoding the “code”
 - Codes  Data B
- Good Auto-Encoder:
 - Data B is very **close** to Data A

Linear Auto-encoder



Deep Auto-Encoder

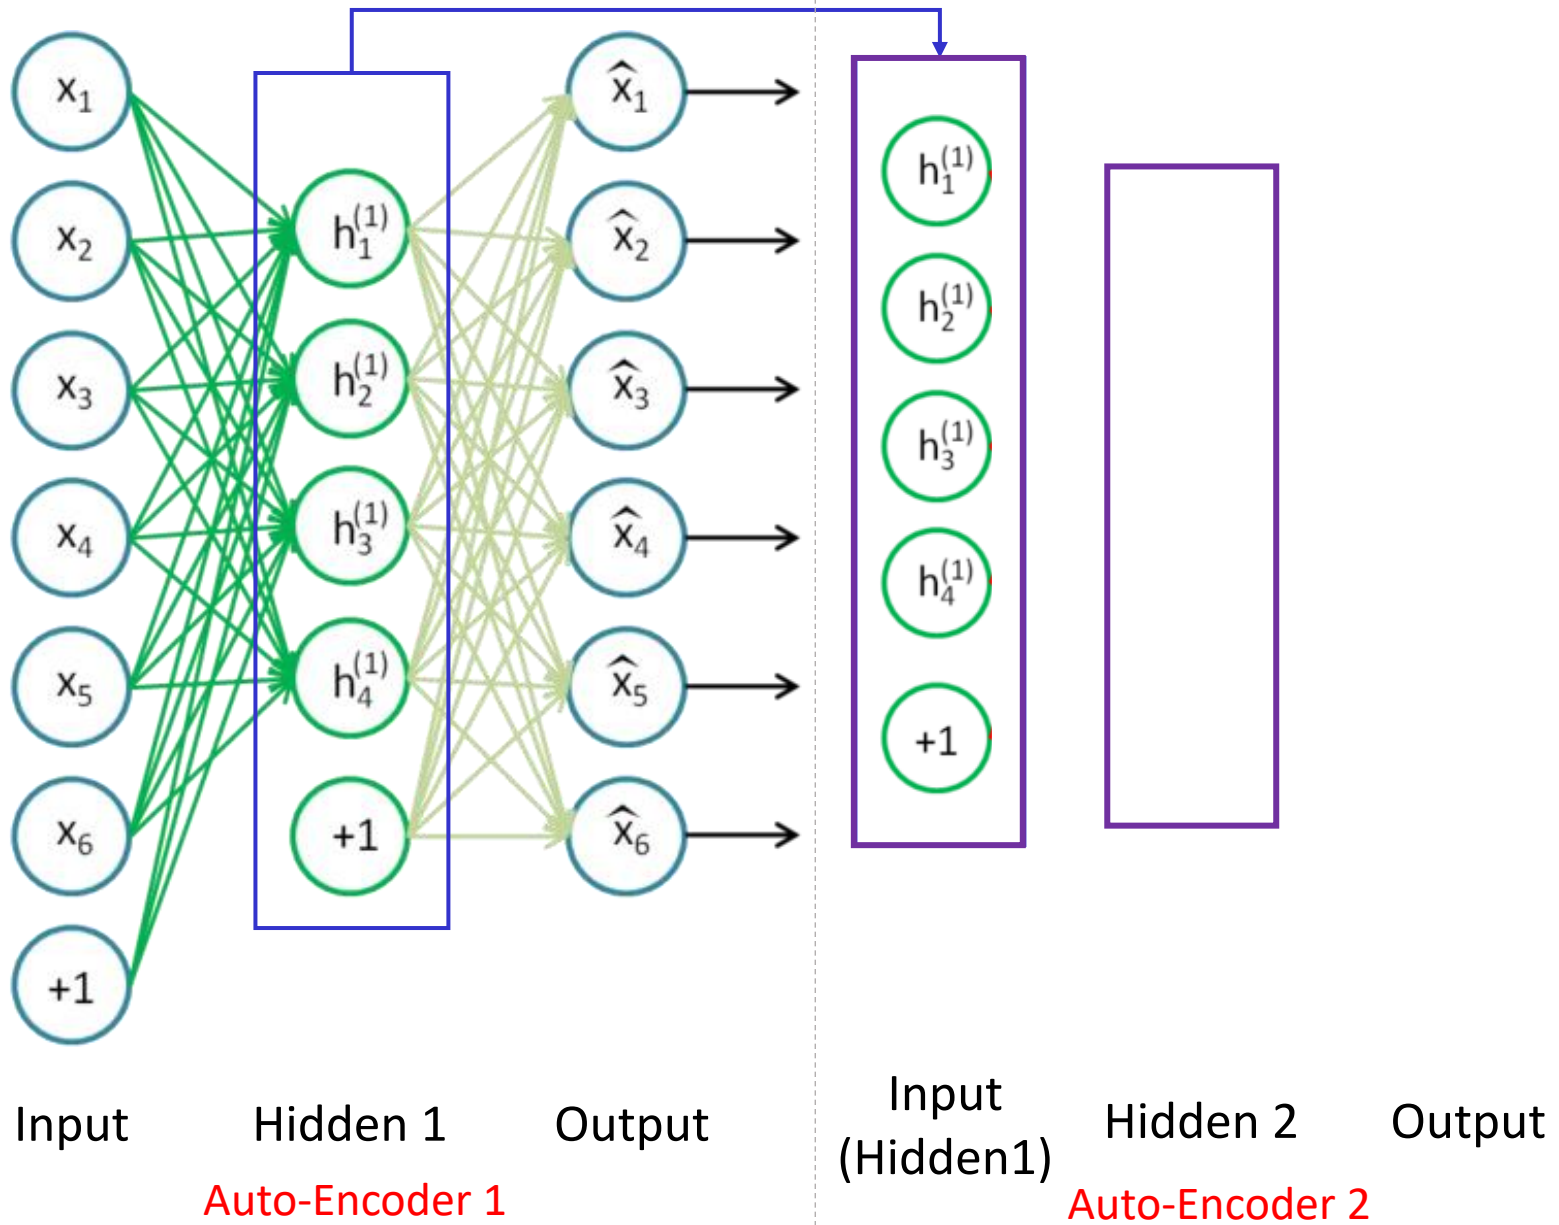
- Of course, the Auto-Encoder can be very deep
 - Not necessary to be symmetric




Stacked Auto-Encoder

- Stacked Auto-Encoder can be used to pre-train the deep network layer by layer
- Two steps
 - Unsupervised pre-training for feature layers
 - Start with the 1st hidden layer, use Auto-Encoder to train this layer to make its inputs and outputs consistent
 - and then use the “code” in the 1st hidden layer as the inputs of the 2nd hidden layer, and repeat...
 - Supervised fine-tuning for classification
 - Compose these pre-trained hidden layers
 - add an output layer at the last, and train the output layer or fine-tune the whole network

Deep Auto-Encoder



Trained Deep Neural Network



Hidden 1 Hidden 2

Convolutional NN

Convolutional Neural Network

- CNN is widely used for image recognition
- It is an end-to-end recognition system
 - A non-linear map that takes raw pixels directly to labels
- Contains the following layers with flexible order and repetitions
 - Convolution layer
 - Activation layer (ReLU: $\max\{0, x\}$)
 - Pooling layer

What is Convolution?




东华帝君




模糊和锐化后的东华帝君

Structure




Input image



Convolution in Neural Networks

- Given an input matrix (e.g. an image)
- Use a small matrix (called **filter** or **kernel**) to screening the input at every position of the input matrix
- Put the convolution results at corresponding positions



Sometimes, we add ReLU activation layer after the outputs

$$output' = \max\{0, output\}$$

Convolutions Visualization


1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Convolutions Visualization (cont.)



Kernel Numbers

Convolution




Pooling

Convolution

Pooling

Fully
ConnectedFully
Connected


Output Predictions



Decide the number
of outputs

Why Convolutions?

- A local in a picture is more important
 - Sparse connections
 - Less computing burden



Why Convolutions?

- A local in a picture is more important
 - Sparse connections
 - Less computing burden
- Position invariance
 - A dog is a dog no matter where he is in the picture
 - Share convolution kernel
 - Share weights

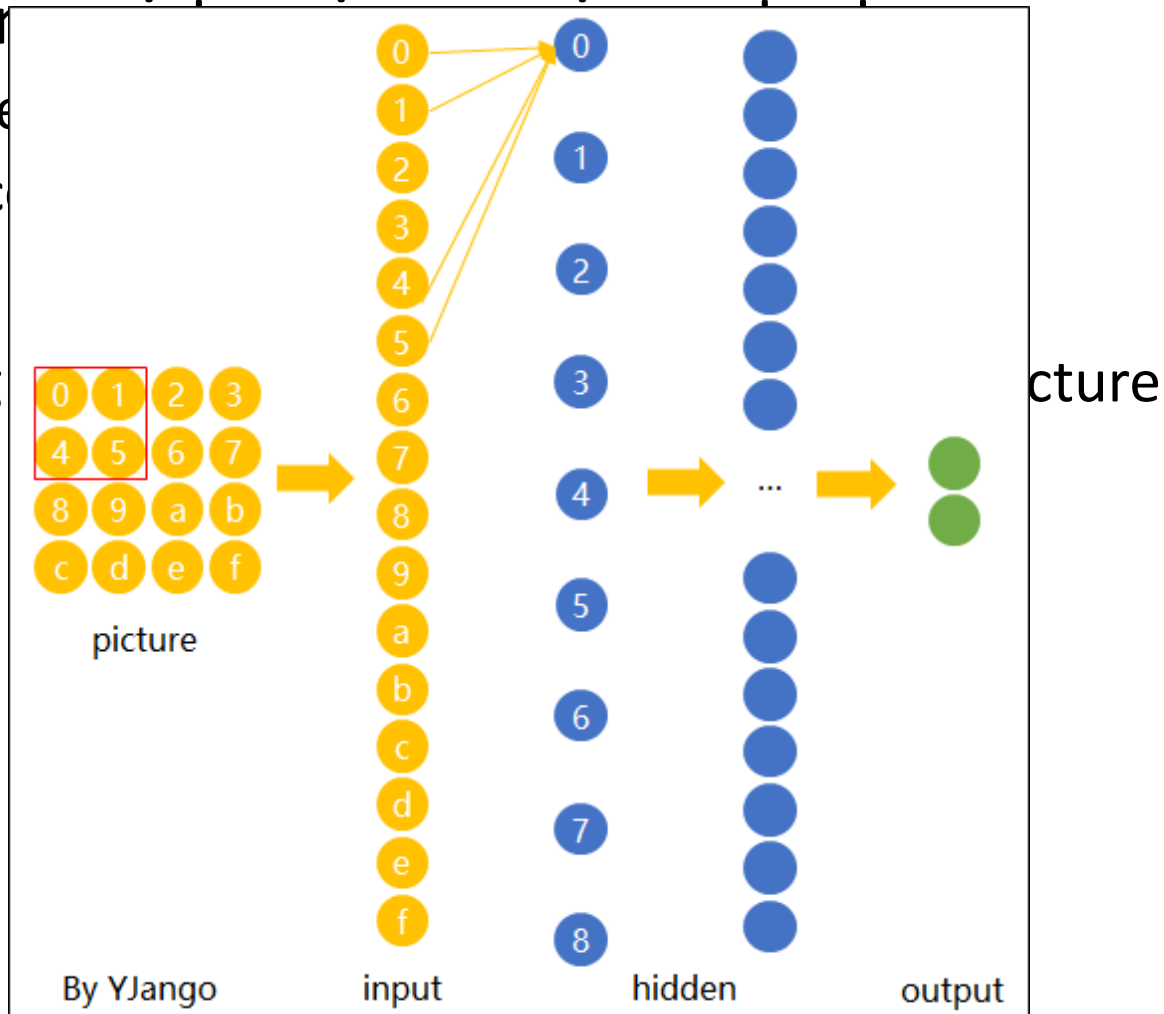
Why Convolutions?

- A local in

- Sparse
- Less c

- Position

- A dog
- Share
- Share



Important Parameters

- Kernel size
 - The dimension of kernel matrix
- Stride
 - The distance that the filter is moved in each step

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature


Kernel Size:
Stride:

Important Parameters

- Kernel size
 - The dimension of kernel matrix
- Stride
 - The distance that the filter is moved in each step
- Pad
 - Add numbers (usually 0) around the input data


Padding

Without padding



$7 \times 7 \Rightarrow 5 \times 5$

With padding




$5 \times 5 \Rightarrow 5 \times 5$

Parameters about Convolution Layer


- Kernel size: k
 - The dimension of kernel matrix is $k \times k$
- Stride: s
 - The distance s that the filter is moved in each step
- Pad: p
 - Add p round of numbers (usually 0) around input data
- Output size calculator:
 - Input size: $w \times h$
 - Output size:
$$w' = \frac{w+2p-k}{s} \qquad h' = \frac{h+2p-k}{s}$$

Pooling Layer


- Make the representations denser and more manageable
- Operate over each activation map independently



Pooling

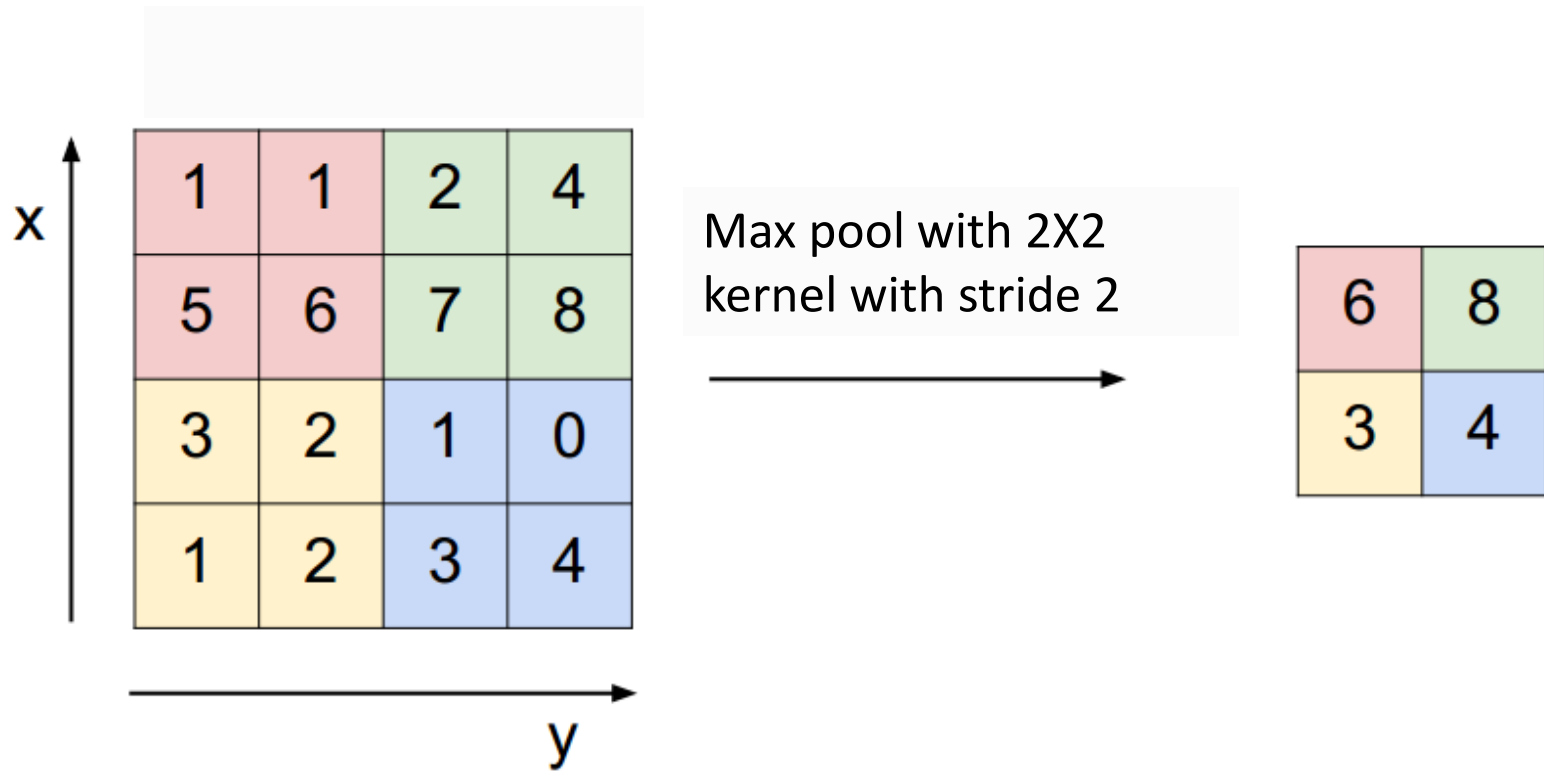


Convolved
feature

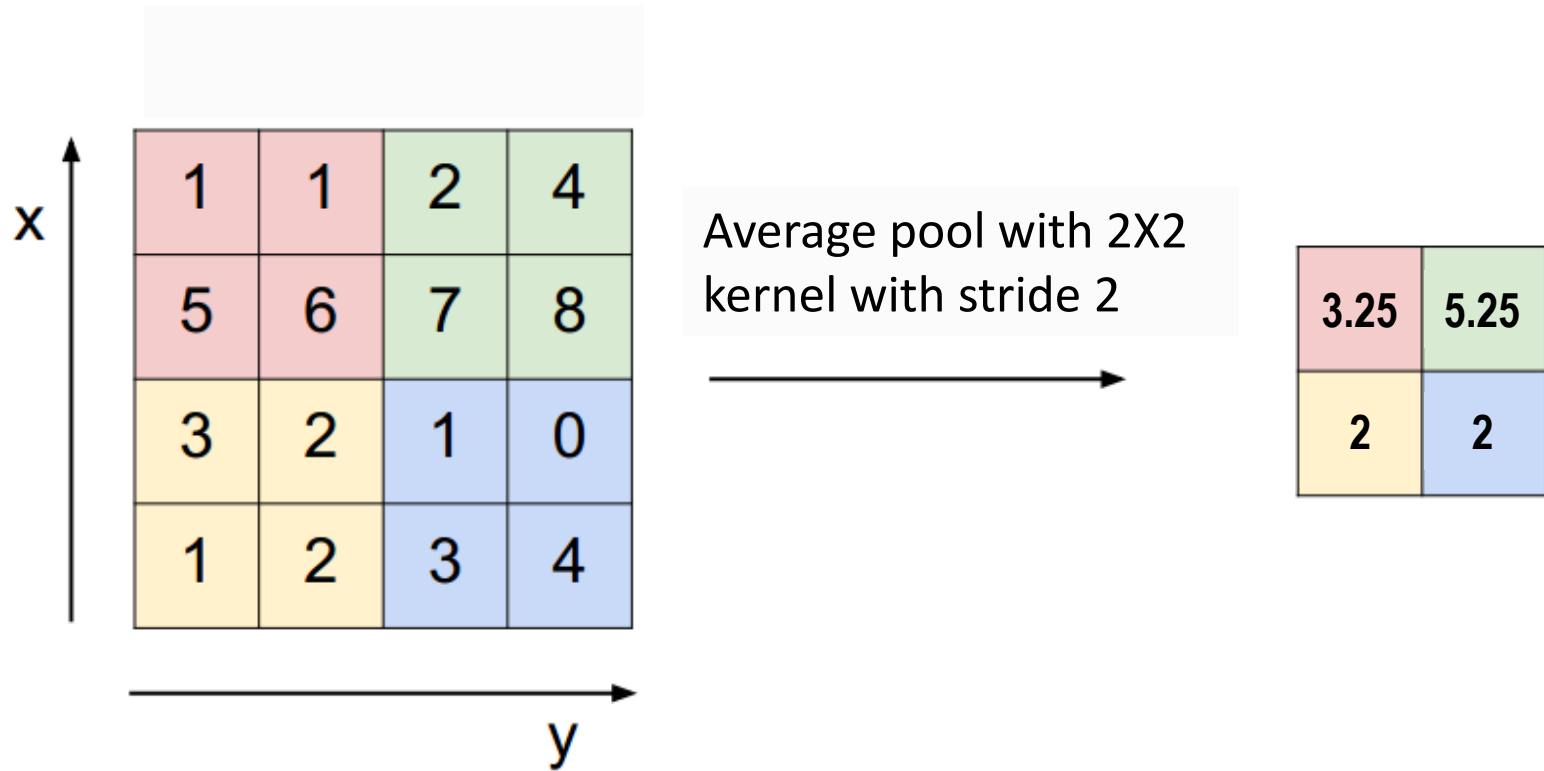


Pooled
feature

Max Pool



Average Pool



Recurrent NN

Motivation of RNN


- In traditional NN
 - Assume all inputs and outputs are independent of each other
 - Input and output length are fixed
- But this might be bad for some tasks
 - Predict next word in a sentence
 - “Context”: You better know which words came before it
- Recurrent
 - Perform the same task for every element of a sequence, with the output being dependent on the previous computations
- They have a “memory” which captures information about what has been calculated so far

Two-layer
feedforward
network



x : input vector
 o : output vector
 s : hidden state vector
 U : layer 1 param. matrix
 V : layer 2 param. Matrix
 f : tanh or ReLU
 $s = f(Ux), o = f(Vs)$

Add time-dependency of the hidden state s




RNN

- x_t is the input at time t
- s_t is the hidden state at time t
 - It is the “memory” of the network
 - Is calculated based on previous hidden state and the input at the current step

$$s_t = f(Ux_t + Ws_{t-1})$$

- o_t is the output at time t

E.g. If we want to predict the next word in a sentence, o_t is a vector of probabilities over certain vocabulary




RNN Features

- s_t is the “memory” of the network
- o_t is based on the memory at t
- RNN share weights U and W
 - Reduce computation complexity
- The output at each time step might be unnecessary
 - E.g. When predicting the sentiment of a sentence we may only care about the final output, not the sentiment after each word
- The input at each time step might be unnecessary
- Most important feature:

The hidden state captures some information about a sequence

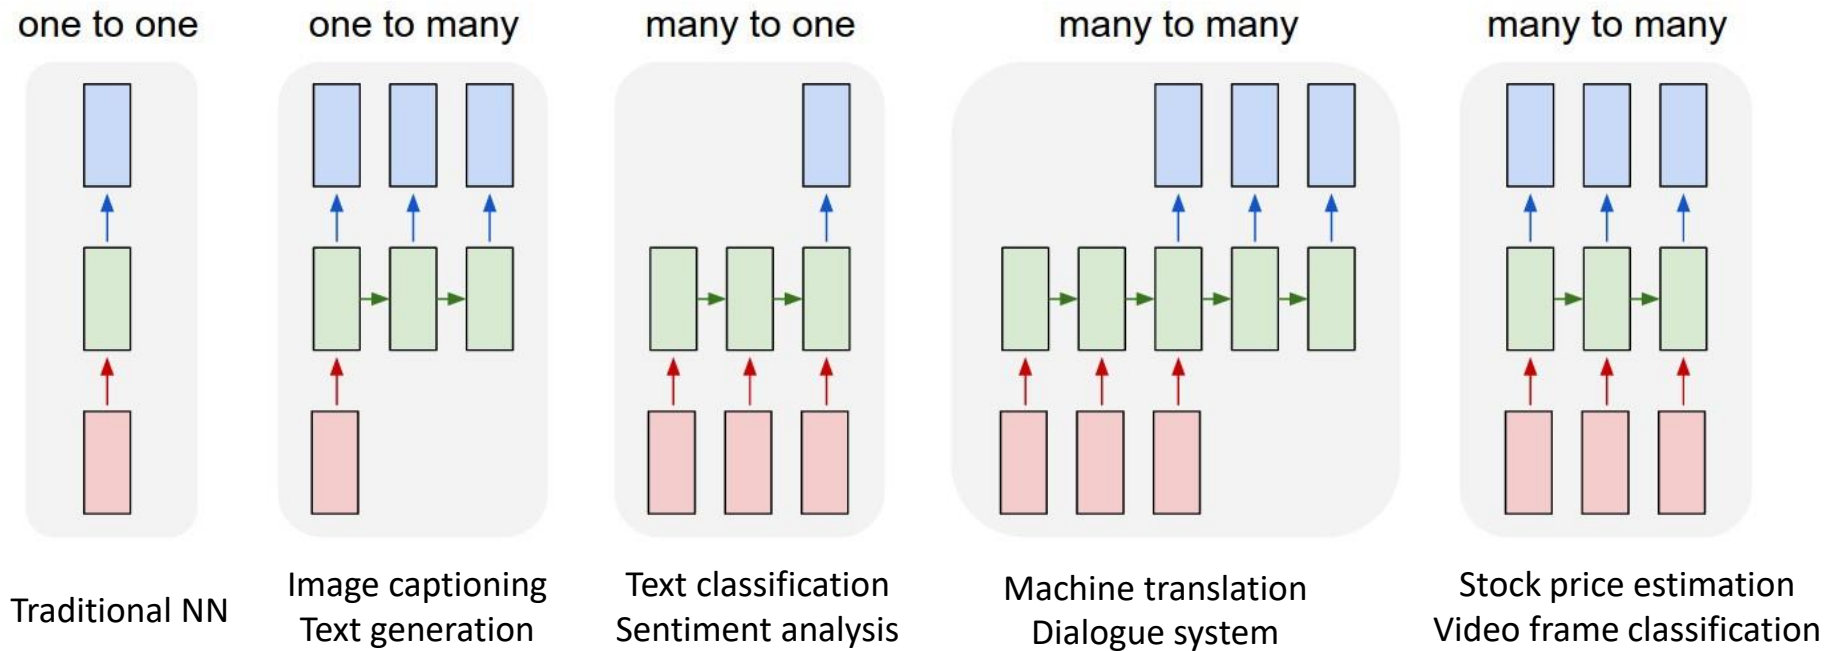
Strategy and Algorithm

- Strategy: minimize the cross entropy
 - E.g., $\hat{y}_t = \text{softmax}(o_t)$, the prediction
 - y_t is the correct word at time t
 - Loss: $-\sum_t y_t \log \hat{y}_t$
- Algorithm: Backpropagation Through Time (BPTT)
 - E.g., in order to calculate the gradient at $t = 4$, we would need to backpropagate 3 steps and sum up the gradients



Different RNN

- Different architecture for various tasks



- Strongly recommend Andrej Karpathy's blog
 - <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Summary

- Universal Approximation: two-layer neural networks can approximate any functions
- Backpropagation is the most important training scheme for multi-layer neural networks so far
- Deep learning, i.e. deep architecture of NN trained with big data, works incredibly well
- Neural networks built with other machine learning models achieve further success

Lecture 9 Wrap-up

- ✓ Deep Learning
- ✓ Deep Auto Encoder
- ✓ Convolutional NN
- ✓ Recurrent NN

Next Lecture

- Supervised learning
 - Linear regression
 - Logistic regression
 - SVM and kernel
 - Tree models
- Deep learning
 - Neural networks
 - Convolutional NN
 - Recurrent NN
- Unsupervised learning
 - Clustering
 - PCA
 - EM
- Reinforcement learning
 - MDP
 - ADP
 - Deep Q-Network



Questions?

Shan Wang (王杉)

<https://wangshan731.github.io/>