

L14: Dimension Reduction

Shan Wang

Lingnan College, Sun Yat-sen University




2020 Data Mining and Machine Learning LN3119

<https://wangshan731.github.io/DM-ML/>



Last lecture

- Unsupervised Learning
- Clustering
 - Hierarchical clustering
 - k -means clustering
- Applications: Netflix



Course Outline

- Supervised learning

- Linear regression
- Logistic regression
- SVM and kernel
- Tree models

- Deep learning

- Neural networks
- Convolutional NN
- Recurrent NN

- Unsupervised learning

- Clustering
- PCA (Dimension Reduction)
- EM

- Reinforcement learning

- MDP
- ADP
- Deep Q-Network

This lecture


- Dimension Reduction
 - Motivation
 - PCA
 - SVD
 - Autoencoder revisit
 - PCA lab
- EM method

Motivation

- Suppose we want to predict the health condition of some students, and the features for the students includes:
 - Weight in kilogram
 - Height in inch
 - Height in cm
 - Hours of sports per day
 - Favorite color
 - Favorite color
- Some features are **irrelevant**, e.g. favorite color and scores in math
- Some features are **redundant**, e.g. height in inch and cm

High dimensional data

- In the era of big data, the dimensionality increases dramatically
 - E.g. there are many features for the electroencephalogram data



- It becomes very important to reduce the dimensionality, or select the most important features, or find the most representative features

Principal Components Analysis

- Principal components analysis (PCA) is a technique that can be used to simplify a dataset
- It is usually a **linear** transformation that chooses a new coordinate system for the data set such that
 - **greatest** variance by any projection of the dataset comes to lie on the **first** axis (then called the first principal component)
 - the second greatest variance on the second axis, and so on
- PCA can be used for reducing dimensionality by eliminating the later principal components

Example

- Consider the following 3D points

1	2	4	3	5	6
2	4	8	6	10	12
3	6	12	9	15	8

- If each component is stored in a byte, we need $18 = 3 \times 6$ bytes

Example (cont.)

- Looking closer, we can see that all the points are related geometrically
 - they are all in the same direction, scaled by a factor:

<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3	$= 1 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3	<table><tr><td>4</td></tr><tr><td>8</td></tr><tr><td>12</td></tr></table>	4	8	12	$= 4 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3	<table><tr><td>5</td></tr><tr><td>10</td></tr><tr><td>15</td></tr></table>	5	10	15	$= 5 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3
1																										
2																										
3																										
1																										
2																										
3																										
4																										
8																										
12																										
1																										
2																										
3																										
5																										
10																										
15																										
1																										
2																										
3																										
<table><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr></table>	2	4	6	$= 2 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3	<table><tr><td>3</td></tr><tr><td>6</td></tr><tr><td>9</td></tr></table>	3	6	9	$= 3 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3	<table><tr><td>6</td></tr><tr><td>12</td></tr><tr><td>8</td></tr></table>	6	12	8	$= 6 \times$	<table><tr><td>1</td></tr><tr><td>2</td></tr><tr><td>3</td></tr></table>	1	2	3
2																										
4																										
6																										
1																										
2																										
3																										
3																										
6																										
9																										
1																										
2																										
3																										
6																										
12																										
8																										
1																										
2																										
3																										

Example (cont.)

1
2
3

 $= 1 \times$

1
2
3

4
8
12

 $= 4 \times$

1
2
3

5
10
15

 $= 5 \times$

1
2
3

2
4
6

 $= 2 \times$

1
2
3

3
6
9

 $= 3 \times$

1
2
3

6
12
8


 $= 6 \times$

1
2
3

- They can be stored using only 9 bytes (50% savings!):
 - Store one direction (3 bytes) + the multiplying constants (6 bytes)

Geometrical interpretation


- View points in 3D space



- In this example, all the points happen to lie on one line
 - a 1D subspace of the original 3D space

Geometrical interpretation


- Consider a new coordinate system where the first axis is along the direction of the line



- In the new coordinate system, every point has only one non-zero coordinate
 - we only need to store the direction of the line (a 3 bytes point) and the nonzero coordinates for each point (6 bytes)

Back to PCA

- Given a set of points, how can we know if they can be compressed similarly to the previous example?
 - We can look into the **correlation** between the points by the tool of **PCA**




From example to theory

- In previous example, PCA rebuilds the coordination system for the data by selecting
 - the direction with **largest variance** as the **first** new base direction
 - the direction with the **second largest variance** as the **second** new base direction
 - and so on
- Then how can we find the direction with largest variance?
 - By the **eigenvector** for the **covariance matrix** of the data

Review – Variance

- Variance is the expectation of the squared deviation of a random variable from its mean
 - Informally, it measures how far a set of (random) numbers are spread out from their average value




Review – Covariance


- Covariance is a measure of the joint variability of two random variables
 - If the greater values of one variable mainly correspond with the greater values of the other variable, (i.e., the variables tend to show similar behavior), the covariance is **positive**
 - E.g. as the number of hours studied increases, the marks in that subject increase
 - In the opposite case, the covariance is **negative**
 - The **sign** of the covariance therefore shows the **tendency** in the **linear relationship** between the variables
 - The magnitude of the covariance is not easy to interpret. The **normalized** version of the covariance, the **correlation coefficient**, however, shows by its **magnitude** the **strength of the linear relation**

Review – Covariance (cont.)


- $\text{covariance}(X, Y) = \frac{\sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{N-1}$




Covariance is -5.4




Covariance is -3




Covariance is 0




Covariance is 2




Covariance is 4.5



PCA



PCA



- PCA tries to identify the subspace in which the data approximately lies in
- PCA uses an **orthogonal transformation** on the coordinate system to convert a set of observations of possibly correlated variables into a set of values of linearly **uncorrelated** variables called principal components
 - The number of principal components is less than or equal to $\min\{d, N\}$

Covariance matrix

- Suppose there are 3 dimensions, denoted as X, Y, Z . The covariance matrix is

$$COV = \begin{bmatrix} COV(X, X) & COV(X, Y) & COV(X, Z) \\ COV(Y, X) & COV(Y, Y) & COV(Y, Z) \\ COV(Z, X) & COV(Z, Y) & COV(Z, Z) \end{bmatrix}$$

- Note the **diagonal** is the covariance of each dimension with respect to itself, which is just the **variance** of each random variable
- Also $COV(X, Y) = COV(Y, X)$
 - hence matrix is **symmetric** about the diagonal
- d -dimensional data will result in a $d \times d$ covariance matrix



Covariance in the covariance matrix

- Diagonal, or the **variance**, measures the deviation from the mean for data points in one dimension
- **Covariance** measures how one dimension random variable varies w.r.t. another, or if there is some linear relationship among them

Data processing

- Given the dataset $D = \{x^{(i)}\}_{i=1}^N$
- Let $\bar{x} = \frac{1}{N} \sum_{i=1}^N x^{(i)}$

- $$X = \begin{bmatrix} (x^{(1)} - \bar{x})' \\ \vdots \\ (x^{(i)} - \bar{x})' \\ \vdots \\ (x^{(N)} - \bar{x})' \end{bmatrix}$$



- Move the center of the data set to 0

Data processing (cont.)

- $Q = X'X =$

$$\begin{bmatrix} (x^{(1)} - \bar{x})' & \dots & (x^{(i)} - \bar{x})' & (x^{(N)} - \bar{x})' \end{bmatrix} \begin{bmatrix} (x^{(1)} - \bar{x})' \\ \vdots \\ (x^{(i)} - \bar{x})' \\ (x^{(N)} - \bar{x})' \end{bmatrix}$$

- Q is square with d dimension
- Q is symmetric
- Q is the **covariance** matrix [aka scatter matrix]
- Q can be very large (in vision, d is often the number of pixels in an image!)
 - or a 256×256 image, $d = 65536!!$
 - Don't want to explicitly compute Q

PCA

- By finding the **eigenvalues** and **eigenvectors** of the covariance matrix, we find that **the eigenvectors with the largest eigenvalues** correspond to the dimensions that have the **strongest variation** in the dataset
- This is the principal component
- Application:
 - face recognition, image compression
 - finding patterns in data of high dimension

The diagram shows the equation $\mathbf{A}\mathbf{X} = \lambda\mathbf{X}$ in a large, bold, black serif font. Two callout boxes with blue borders and white backgrounds point to parts of the equation. The top callout box points to the \mathbf{X} in $\mathbf{A}\mathbf{X}$ and contains the text "Eigenvector of Matrix A". The bottom callout box points to the λ and contains the text "Eigenvalue of Matrix A".

$$\mathbf{A}\mathbf{X} = \lambda\mathbf{X}$$

PCA theorem

- Theorem:
- Each $x^{(i)}$ can be written as: $x^{(i)} = \bar{x} + \sum_{j=1}^d g_{ij} e_j$
where e_j are the d eigenvectors of Q with non-zero eigenvalues
- Notes:
 1. The eigenvectors e_1, e_2, \dots, e_d span an eigenspace
 2. e_1, e_2, \dots, e_d are $d \times 1$ orthonormal vectors
(directions in d -Dimensional space)
 3. The scalars g_{ij} are the coordinates of $x^{(i)}$ in the space
$$g_{ij} = \langle x^{(i)} - \bar{x}, e_j \rangle$$

Using PCA to compress data

- Expressing x in terms of e_1, e_2, \dots, e_d doesn't change the size of the data
- Sort the eigenvectors e_j according to their eigenvalue

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$$

- Assume $\lambda_j \approx 0$ if $j > k$. Then

$$x^{(i)} \approx \bar{x} + \sum_{j=1}^k g_{ij} e_j$$

Bonus question:

$$\lambda_1 + \lambda_2 + \dots + \lambda_d = ?$$

How to solve eigenvalues and eigenvectors?

- If $(A - \lambda I)x = 0$ has a nonzero solution for λ , then $A - \lambda I$ is not invertible. Then the determinant of $A - \lambda I$ must be zero
- λ is an eigenvalue of A if and only if $A - \lambda I$ is singular:

$$\det(A - \lambda I) = 0$$

The diagram shows the equation $\mathbf{Ax} = \lambda \mathbf{x}$. A callout box labeled "Eigenvector of Matrix A" points to the variable \mathbf{x} . Another callout box labeled "Eigenvalue of Matrix A" points to the scalar λ .

Example

- Find the eigenvalues and eigenvectors of $A = \begin{bmatrix} 1 & 2 \\ 2 & 4 \end{bmatrix}$
- Need to solve $\det(A - \lambda I) = 0$
- $A - \lambda I = \begin{bmatrix} 1 - \lambda & 2 \\ 2 & 4 - \lambda \end{bmatrix}$
- $\det(A - \lambda I) = (1 - \lambda)(4 - \lambda) - 2 \times 2 = \lambda^2 - 5\lambda$
- $\det(A - \lambda I) = 0$ would imply $\lambda = 0$ and $\lambda = 5$

Example (cont.)

- then solve $(A - \lambda I)x = 0$ to get the eigenvectors for each of the eigenvalues
- $Ax = 0$ has a nonzero solution of $\begin{bmatrix} 2 \\ -1 \end{bmatrix}$ or $\begin{bmatrix} 2/\sqrt{5} \\ -1/\sqrt{5} \end{bmatrix}$
- $(A - 5I)x = 0$ has a nonzero solution of $\begin{bmatrix} 1 \\ 2 \end{bmatrix}$ or $\begin{bmatrix} 1/\sqrt{5} \\ 2/\sqrt{5} \end{bmatrix}$
- Note: Eigenvectors for different eigenvalues are **orthogonal**

How about non-squared matrix A ?

SVD

Singular Value Decomposition

SVD

- Singular Value Decomposition (SVD) is a factorization method of matrix. It states that any $m \times n$ matrix A can be written as the product of 3 matrices:

$$A = USV^T$$

- Where:
 - U is $m \times m$ and its columns are orthonormal **eigenvectors** of AA^T
 - V is $n \times n$ and its columns are orthonormal eigenvectors of $A^T A$
 - S is $m \times n$ is a diagonal matrix with r elements equal to the root of the positive eigenvalues of AA^T or $A^T A$ (both matrices have the same positive eigenvalues anyway)

In full matrix form

$$A_{m \times n} = U_{m \times m} S_{m \times n} V_{n \times n}^T$$

$$\begin{array}{c} \text{A} \\ \left(\begin{array}{ccc} x_{11} & x_{12} & x_{1n} \\ \vdots & \ddots & \vdots \\ x_{m1} & & x_{mn} \end{array} \right) \\ m \times n \end{array} = \begin{array}{c} \text{U} \\ \left(\begin{array}{ccc} u_{11} & & u_{m1} \\ & \ddots & \\ u_{1m} & & u_{mm} \end{array} \right) \\ m \times m \end{array} \begin{array}{c} \text{S} \\ \left(\begin{array}{ccc} \sigma_1 & & 0 \\ & \ddots & \\ 0 & \sigma_r & \\ & & \ddots & \\ & & 0 & \end{array} \right) \\ m \times n \end{array} \begin{array}{c} \text{V}^T \\ \left(\begin{array}{ccc} v_{11} & & v_{1n} \\ & \ddots & \\ v_{n1} & & v_{nn} \end{array} \right) \\ n \times n \end{array}$$

Example

- Let's assume:

$$A = \begin{pmatrix} 3 & 2 & 2 \\ 2 & 3 & -2 \end{pmatrix}$$

- We can have:

$$AA^T = \begin{pmatrix} 17 & 8 \\ 8 & 17 \end{pmatrix} \quad A^T A = \begin{pmatrix} 13 & 12 & 2 \\ 12 & 13 & -2 \\ 2 & -2 & 8 \end{pmatrix}$$

Example (cont.)

- Compute U and V respectively:

$$AA^T = \begin{pmatrix} 17 & 8 \\ 8 & 17 \end{pmatrix}$$

eigenvalues: $\lambda_1 = 25, \lambda_2 = 9$

eigenvectors

$$u_1 = \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} \quad u_2 = \begin{pmatrix} 1/\sqrt{2} \\ -1/\sqrt{2} \end{pmatrix}$$

$$A^T A = \begin{pmatrix} 13 & 12 & 2 \\ 12 & 13 & -2 \\ 2 & -2 & 8 \end{pmatrix}$$

eigenvalues: $\lambda_1 = 25, \lambda_2 = 9, \lambda_3 = 0$

eigenvectors

$$v_1 = \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \\ 0 \end{pmatrix} \quad v_2 = \begin{pmatrix} 1/\sqrt{18} \\ -1/\sqrt{18} \\ 4/\sqrt{18} \end{pmatrix} \quad v_3 = \begin{pmatrix} 2/3 \\ -2/3 \\ -1/3 \end{pmatrix}$$

Example (cont.)



- Finally, we have:

$$A = USV^T = \begin{pmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{pmatrix} \begin{pmatrix} 5 & 0 & 0 \\ 0 & 3 & 0 \end{pmatrix} \begin{pmatrix} 1/\sqrt{2} & 1/\sqrt{2} & 0 \\ 1/\sqrt{18} & -1/\sqrt{18} & 4/\sqrt{18} \\ 2/3 & -2/3 & -1/3 \end{pmatrix}$$

Insight



$$A = U S V^T = \underbrace{\sigma_1}_{\text{more significant}} \underbrace{u_1}_{m \times 1} \underbrace{v_1^T}_{1 \times n} + \dots + \underbrace{\sigma_r}_{\text{less significant}} \underbrace{u_r}_{m \times 1} \underbrace{v_r^T}_{1 \times n}$$

$$S = \sigma_1 u_1 v_1^T + \cancel{\sigma_2 u_2 v_2^T}$$

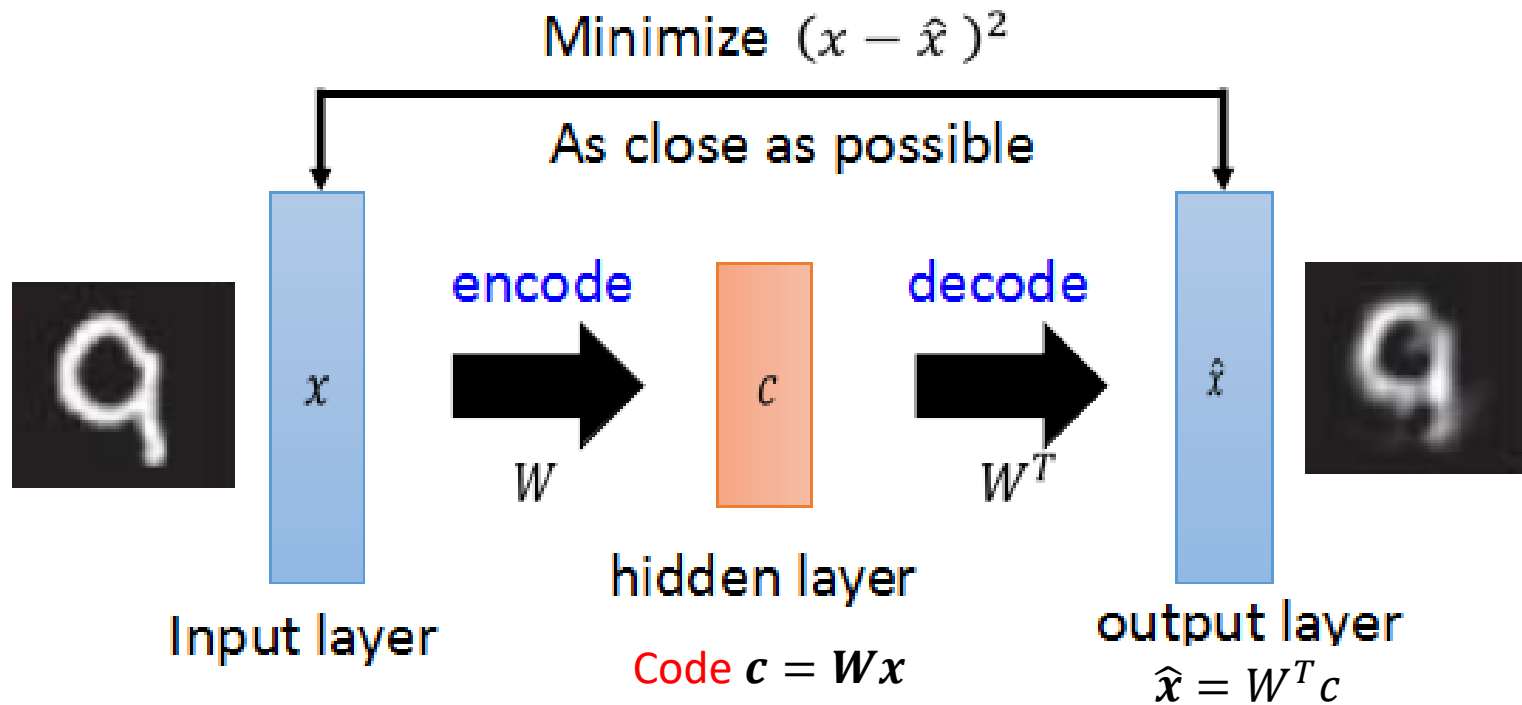


Autoencoder revisit

Auto-Encoder

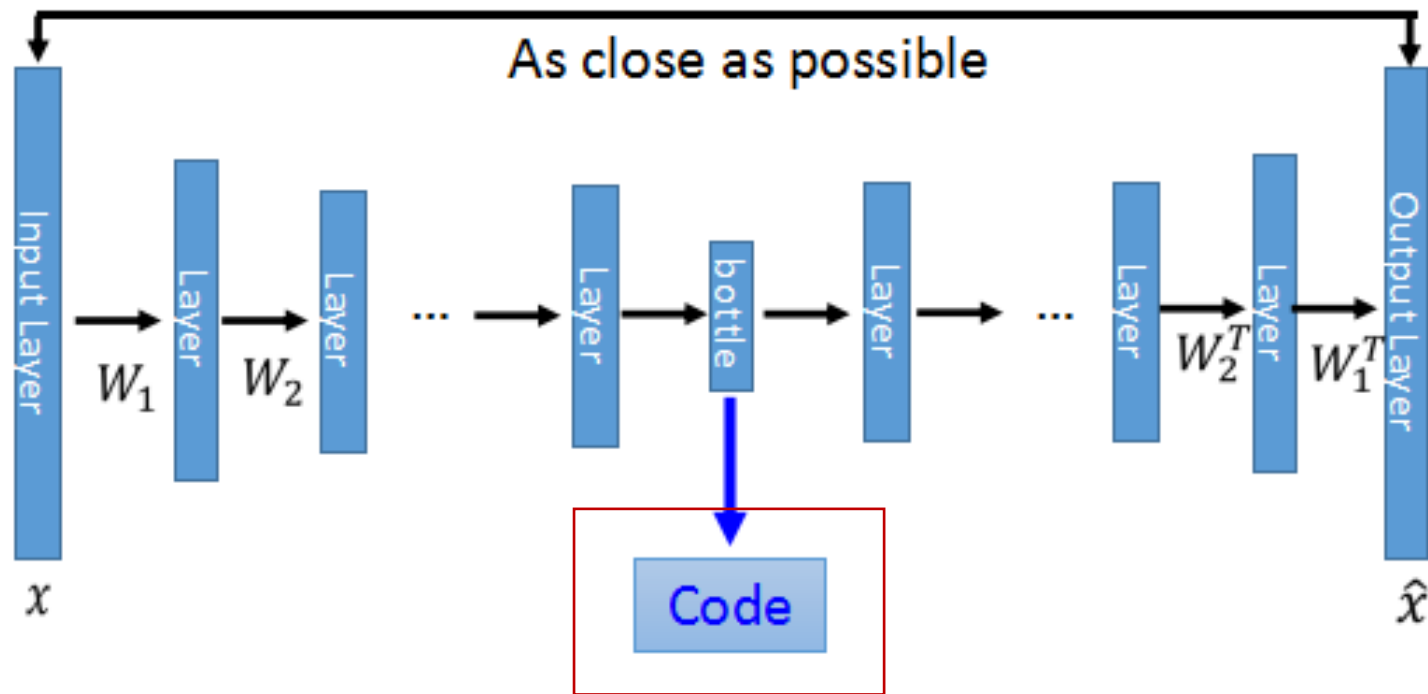
- Auto-Encoder can be used for dimensionality reduction.
- An auto-encoder contains an **encoder** and a **decoder**
- The encoder plays the role of coding the original data
 - Data A  Codes
- The decoder plays the role of decoding the “code”
 - Codes  Data B
- Good Auto-Encoder:
 - Data B is very **close** to Data A

Linear Auto-encoder



Deep Auto-Encoder

- Of course, the Auto-Encoder can be very deep
 - Not necessary to be symmetric



Dimension Reduced

Let's get our hands dirty!

Data Analysis

Let's do some basic data analysis using our WHO data.

Hide

WHO\$Under15

```
[1] 47.42 21.33 27.42 15.20 47.58 25.96 24.42 20.34 18.95 14.51 22.25 21.62 20.16 30.57 18.99 15.10 16.88 34.4
0 42.95 28.53
[21] 35.23 16.35 33.75 24.56 25.75 13.53 45.66 44.20 31.23 43.88 16.37 30.17 40.07 48.52 21.38 17.95 28.03 42.1
7 42.37 30.61
[41] 23.94 41.68 14.08 16.58 17.16 14.56 21.98 45.11 17.66 33.72 25.96 30.53 30.29 31.25 30.63 38.95 43.10 15.6
9 43.29 28.88
[61] 16.42 18.26 38.49 45.90 17.62 13.17 38.59 14.60 26.96 40.80 42.46 41.55 36.77 35.35 35.72 14.62 20.71 29.4
3 29.27 23.68
[81] 40.51 23.54 27.53 14.84 27.78 13.13 34.13 25.46 42.37 30.10 24.90 30.21 35.61 14.57 21.64 36.75 43.05 29.4
5 15.13 17.46
[101] 42.72 45.44 26.65 29.83 47.14 14.98 30.10 40.22 20.17 29.82 35.81 18.26 27.85 19.81 27.85 45.38 25.28 36.5
9 38.10 35.58
[121] 17.21 20.26 33.37 49.99 44.23 30.61 18.64 24.19 34.31 30.10 28.65 38.37 32.78 29.18 34.53 14.91 14.92 13.2
8 15.25 16.52
[141] 15.05 15.45 43.56 25.96 24.31 25.70 37.88 14.04 41.60 29.69 43.54 16.45 21.95 41.74 16.48 15.00 14.16 40.3
7 47.35 29.53
[161] 42.28 15.20 25.15 41.48 27.83 38.05 16.71 14.79 35.35 35.75 18.47 16.89 46.33 41.89 37.33 20.73 23.22 26.0
0 28.65 30.61
[181] 48.54 14.18 14.41 17.54 44.85 19.63 22.05 28.90 37.37 28.84 22.87 40.72 46.73 40.34
```


WHO\$Country[which.min(WHO\$Under15)]

```
[1] Japan
194 levels: Afghanistan Albania Algeria Andorra Angola Antigua and Barbuda Argentina Armenia Australia Austria
... Zimbabwe
```

Let's create some plots for exploratory data analysis (EDA). First, let's create a basic scatterplot of GNI versus FertilityRate.

Hide

plot(WHO\$GNI, WHO\$FertilityRate)



EM Method

Problem

- Given the training dataset

$$D = \{x_i\}_{i=1,2,\dots,N}$$

Let the machine learn the data underlying patterns

- In last lecture, we have learned Clustering
- In EM, we assume latent variables

$$Z \rightarrow \mathbf{x}$$

- We wish to fit the parameters of a model $p(\mathbf{x}, z)$ to the data, where the log-likelihood is

$$\begin{aligned} l(\theta) &= \sum_{i=1}^N \log p(x; \theta) \\ &= \sum_{i=1}^N \log \sum_z p(x, z; \theta) \end{aligned}$$

Method

- Explicitly find the maximum likelihood estimation (MLE) is hard

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^N \log \sum_z p(x^{(i)}, z^{(i)}; \theta)$$

- But given $z^{(i)}$ observed, the MLE is easy

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^N \log p(x^{(i)} | z^{(i)}; \theta)$$

- EM methods give an efficient solution for MLE, by iteratively doing
 - E-step: **construct a (good) lower-bound** of log-likelihood
 - M-step: optimize that lower-bound

How to construct a good lower-bound?

Lower Bound

- For each instance i , let q_i be some distribution of $z^{(i)}$


$$\sum_z q_i(z) = 1, \quad q_i(z) \geq 0$$

- Thus the data log-likelihood


$$\begin{aligned}
 l(\theta) &= \sum_{i=1}^N \log p(x^{(i)}; \theta) = \sum_{i=1}^N \log \sum_{z^{(i)}} p(x^{(i)}, z^{(i)}; \theta) \\
 &= \sum_{i=1}^N \log \sum_{z^{(i)}} q_i(z^{(i)}) \frac{p(x^{(i)}, z^{(i)}; \theta)}{q_i(z^{(i)})} \\
 &\geq \sum_{i=1}^N \sum_{z^{(i)}} q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{q_i(z^{(i)})} \quad \text{Lower bound of } l(\vartheta)
 \end{aligned}$$

Jensen's inequality
 $-\log(x)$ is a convex function

Jensen's Inequality



Jensen's Inequality



Lower bound

$$l(\theta) = \sum_{i=1}^N \log p(x^{(i)}; \theta) \geq \sum_{i=1}^N \sum_{z^{(i)}} q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{q_i(z^{(i)})}$$

- Then what $q_i(z)$ should we choose?
- The posterior distribution

$$q_i(z^{(i)}) = \frac{p(x^{(i)}, z^{(i)}; \theta)}{\sum_z p(x^{(i)}, z; \theta)} = \frac{p(x^{(i)}, z^{(i)}; \theta)}{p(x^{(i)}; \theta)} = p(z^{(i)} | x^{(i)}; \theta)$$

General EM Methods

- Repeat until convergence: {

- (E-step) For each i , set

$$q_i(z^{(i)}) = p(z^{(i)} | x^{(i)}; \theta)$$

- (M-step) Update the parameters

$$\theta = \arg \max_{\theta} \sum_{i=1}^N \sum_{z^{(i)}} q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{q_i(z^{(i)})}$$

}

Lecture 14 Wrap-up

- ✓ Dimension Reduction

 - ✓ Motivation

 - ✓ PCA

 - ✓ SVD

 - ✓ Autoencoder revisit

 - ✓ PCA lab

- ✓ EM method

Next Lecture

- Supervised learning
 - Linear regression
 - Logistic regression
 - SVM and kernel
 - Tree models
- Deep learning
 - Neural networks
 - Convolutional NN
 - Recurrent NN
- Unsupervised learning
 - Clustering
 - PCA (Dimension Reduction)
 - EM
- Reinforcement learning
 - MDP
 - ADP
 - Deep Q-Network



Questions?

Shan Wang (王杉)

<https://wangshan731.github.io/>