

L15: Introduction to Reinforcement Learning

Shan Wang

Lingnan College, Sun Yat-sen University

2020 Data Mining and Machine Learning LN3119

<https://wangshan731.github.io/DM-ML/>



Last lecture


- Dimension Reduction
 - PCA (squared matrix)
 - SVD (general matrix)
 - Auto Encoder

$$\mathbf{Ax} = \lambda \mathbf{x}$$

Eigenvector of Matrix A

Eigenvalue of Matrix A

$$A_{m \times n} = U_{m \times m} S_{m \times n} V_{n \times n}^T$$



- EM Method
 - E-step: **construct a (good) lower-bound** of log-likelihood
 - M-step: optimize that lower-bound

What we have learned so far

- Supervised Learning
 - To perform the desired output given the data and labels
 - e.g., to build a loss function to minimize
- Unsupervised Learning
 - To analyze and make use of the underlying data patterns/structures
 - e.g., to build a log-likelihood function to maximize

Supervised Learning

- Given the training dataset of (data, label) pairs,

$$D = \{(\mathbf{x}_i, y_i)\}_{i=1,2,\dots,N}$$

let the machine learn a function from data to label

$$y_i \simeq f_{\theta}(\mathbf{x}_i)$$

- Learning is referred to as updating the parameter θ
- Learning objective: make the prediction close to the ground truth
 - $f_{\theta}(\mathbf{x}_i)$ is as close to y_i as possible

Unsupervised Learning

- Given the training dataset

$$D = \{(\mathbf{x}_i)\}_{i=1,2,\dots,N}$$

let the machine learn the data underlying patterns

- Sometimes build latent variables

$$z \rightarrow \mathbf{x}$$

- Estimate the probabilistic density function (p.d.f.)

$$p(\mathbf{x}|\theta) = \sum_z q(z)p(\mathbf{x}|z, \theta)$$

- Maximize the likelihood of training data

$$\prod_{i=1}^N p(\mathbf{x}_i|\theta)$$

Two Kinds of Machine Learning

- Prediction
 - Predict the desired output given the data (supervised learning)
 - Generate data instances (unsupervised learning)
 - We mainly covered this category in previous lectures
- Decision Making
 - Take actions based on a particular state in a dynamic environment (reinforcement learning)
 - to transit to new states
 - to receive immediate reward
 - to maximize the accumulative reward over time
 - Learning from interaction

Machine Learning Categories

- Supervised Learning

- To perform the desired output given the data and labels

$$p(y|\mathbf{x})$$

- Unsupervised Learning

- To analyze and make use of the underlying data patterns/structures

$$p(\mathbf{x})$$

- Reinforcement Learning

- To learn a policy of taking actions in a dynamic environment and acquire rewards

$$\pi(a|\mathbf{x})$$

Course Outline

- Supervised learning

- Linear regression
- Logistic regression
- SVM and kernel
- Tree models

- Deep learning

- Neural networks
- Convolutional NN
- Recurrent NN

- Unsupervised learning

- Clustering
- PCA
- EM

- Reinforcement learning

- MDP
- ADP
- Deep Q-Network


This lecture

- Introduction to Reinforcement Learning
- Model-based Reinforcement Learning
 - MDP
 - Value iteration
 - Policy iteration
- Model-free Reinforcement Learning

Introduction to RL


Reinforcement Learning

- Learning from interaction
 - Given the current situation, what to do next in order to maximize utility?



Reinforcement Learning Definition

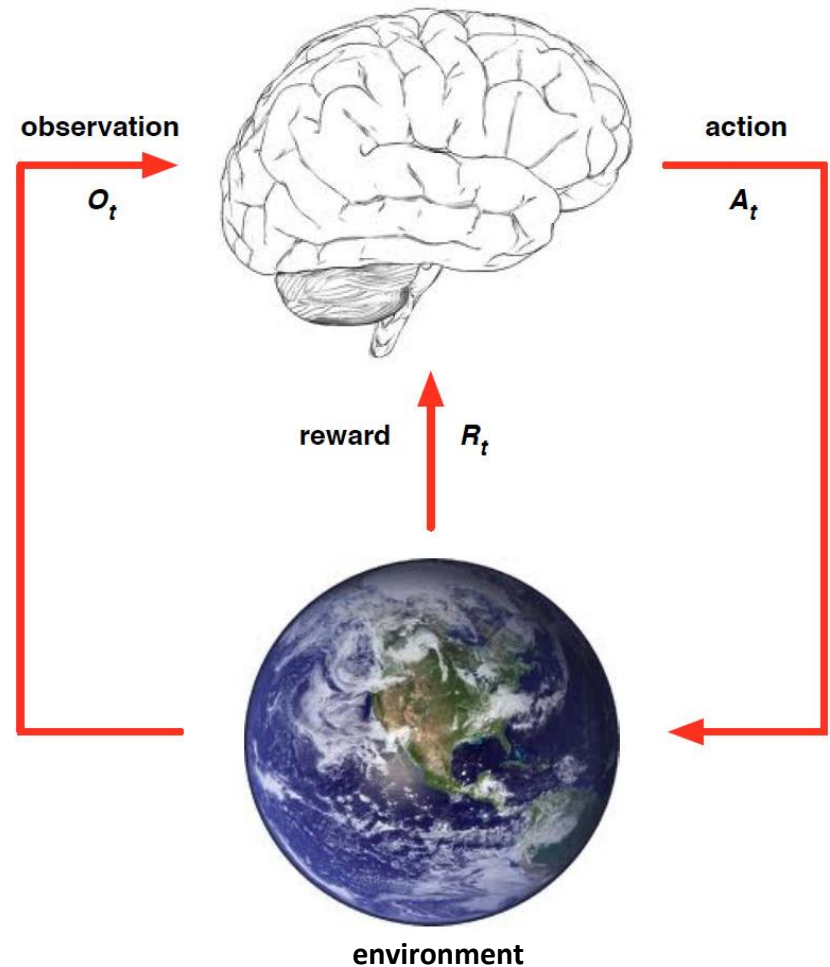
- A computational approach by learning from interaction to achieve a goal



- Three aspects
 - Sensation: sense the state of the environment to some extent
 - Action: able to take actions that affect the state and achieve the goal
 - Goal: maximize the cumulative reward over time

Process of RL

- At step t , the agent
 - Receives observation O_t
 - Receives scalar reward R_t
 - Executes action A_t
- The environment
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits scalar reward R_{t+1}
- t increments at environment step



Elements of RL Systems

- **History** is the sequence of observations, action, rewards

$$H_t = \{O_1, R_1, A_1, O_2, R_2, A_2, \dots, O_t, R_t\}$$

- i.e. all observable variables up to time t
- E.g., all the records of the Go game
- What happens next depends on the history:
 - The agent selects actions
 - The environment selects observations/rewards
- **State** is the information used to determine what happens next (actions, observations, rewards)
- Formally, state is a function of the history

$$S_t = f(H_t)$$

Elements of RL Systems

- **Policy** is the learning agent's way of behaving at a given time

- It is a map from state to action
- Deterministic policy

$$a = \pi(s)$$

- Stochastic policy

$$\pi(s|a) = P(A_t = a | S_t = s)$$

Elements of RL Systems

- Reward
 - A scalar defining the goal in an RL problem
 - For immediate sense of what is good
- Value function
 - State value (Value of a state) is a scalar specifying what is good in the long run
 - Value function is a prediction of the cumulative future reward
 - Used to evaluate the goodness/badness of states (given the current policy)

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | S_t = s]$$

Elements of RL Systems

- Reward

- A scalar defining the goal in an RL problem
- For immediate sense of what is good

- Value function

- State value (value of a state) is a scalar specifying what is good in the long run, i.e., the cumulative reward

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | S_t = s]$$

- Action value (value of a action) is a scalar specifying what is a good action at a specific state in the long run

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | S_t = s, A_t = a]$$

Elements of RL Systems

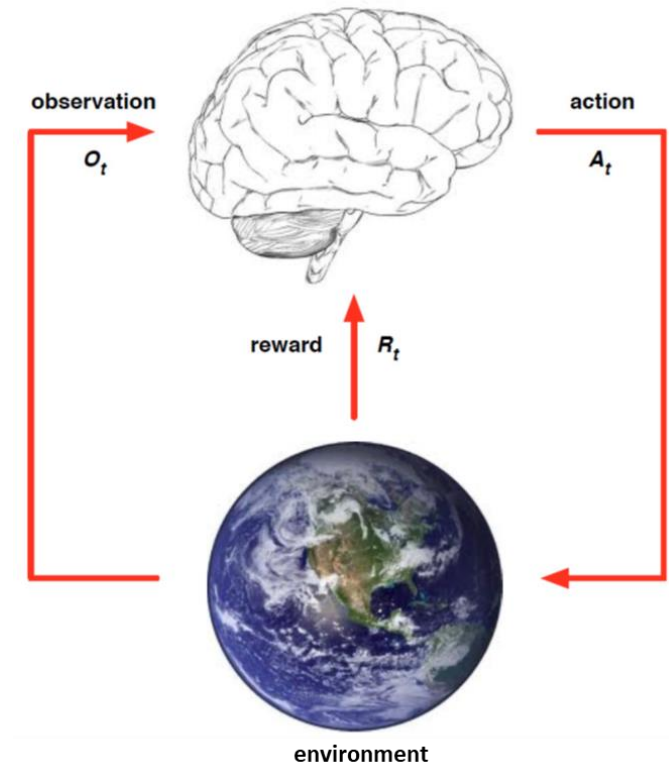
- A Model of the environment that mimics the behavior of the environment

- Predict the next state

$$\mathcal{P}_{sa}(s') = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$$


- Predict the next (immediate) reward

$$\mathcal{R}_s(a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$




Maze Example

- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
 - No move if the action is to the wall




Maze Example

- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step



Maze Example

- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step




Given a policy as shown above

- Arrows represent policy $\pi(s)$ for each state s

Maze Example

- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step



Numbers represent value $v_{\pi}(s)$ of each state s

Another example

- http://www.4399.com/flash/105474_1.htm

Model-based RL

Markov Decision Process

Markov Decision Process

- Markov decision processes (MDPs) provide a mathematical framework for modeling decision making in situations where outcomes are partly random and partly under the control of a decision maker.
- MDPs formally describe an environment for RL
 - where the environment is FULLY observable
 - i.e. the current state completely characterizes the process (Markov property)

Markov Property

- “The future is independent of the past given the present”
- Definition
 - A state s_t is Markov if and only if
$$P[s_{t+1}|s_t] = P[s_{t+1}|s_t, s_{t-1}, \dots, s_1]$$
- Properties
 - The state captures all relevant information from the history
 - Once the state is known, the history may be thrown away
 - i.e. the state is sufficient statistic of the future

Markov Decision Process

- A Markov decision process is a tuple $(S, A, \{P_{sa}\}, \gamma, R)$
 - S is the set of states
 - E.g., location in a maze, or current screen in an Atari game
 - A is the set of actions
 - E.g., move N, E, S, W, or the direction of the joystick and the buttons
 - P_{sa} are the state transition probabilities
 - For each state $s \in S$ and action $a \in A$, P_{sa} is a distribution over the next state in S
 - $\gamma \in [0,1]$ is the discount factor for the future reward
 - $R: S \times A \rightarrow \mathbb{R}$ is the reward function
 - Sometimes the reward is only assigned to state, i.e., irrelative to the action

Markov Decision Process

- The dynamics of an MDP proceeds as
 - Start in a state s_0
 - The agent chooses some action $a_0 \in A$
 - The agent gets the reward $R(s_0, a_0)$
 - MDP randomly transits to some successor state $s_1 \sim P_{s_0 a_0}$
 - This proceeds iteratively

$$\begin{array}{ccccccc} & a_0 & & a_1 & & a_2 & \\ s_0 & \xrightarrow{\quad} & s_1 & \xrightarrow{\quad} & s_2 & \xrightarrow{\quad} & s_3 \dots \\ & R(s_0, a_0) & & R(s_1, a_1) & & R(s_2, a_2) & \end{array}$$

- Until a terminal state S_T or proceeds with no end
- The total payoff of the agent is

$$R(s_0, a_0) + \gamma R(s_1, a_1) + \gamma^2 R(s_2, a_2) + \dots$$

Reward on State Only

- For a large part of cases, reward is only assigned to the state
 - E.g., in maze game, the reward is on the location
 - In game of Go, the reward is only based on the final territory

- The reward function $R(s): S \rightarrow \mathbb{R}$

- MDPs proceed

$$\begin{array}{ccccccc} & a_0 & & a_1 & & a_2 & \\ s_0 & \xrightarrow{\quad} & s_1 & \xrightarrow{\quad} & s_2 & \xrightarrow{\quad} & s_3 \dots \\ & R(s_0) & & R(s_1) & & R(s_2) & \end{array}$$

- cumulative reward (total payoff)

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

MDP Goal and Policy

- The goal is to choose actions over time to maximize the expected cumulative reward


$$\mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

- $\gamma \in [0,1]$ is the discount factor for the future reward, which makes the agent prefer immediate reward to future reward
 - In finance case, today's \$1 is more valuable than \$1 in tomorrow
- Given a particular policy $\pi(s): S \rightarrow A$
 - i.e. take the action $a = \pi(s)$ at state s
- Define the value function for π
$$V^\pi(s) = \mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi]$$
 - i.e. expected cumulative reward given the start state s and taking actions according to π

Bellman Equation for Value Function


- Define the value function for π

$$V^\pi(s) = \mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots | s_0 = s, \pi]$$


$$\gamma V^\pi(s_1)$$

Bellman Equation:

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in \mathcal{S}} P_{s\pi(s)}(s') V^\pi(s')$$




Immediate
Reward



Time
decay



State
transition



Value of
the next
state

Optimal Value Function

- The optimal value function for each state s is best possible sum of discounted rewards that can be attained by any policy

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

- The Bellman's equation for optimal value function

$$V^*(s) = R(s) + \gamma \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

- The optimal policy

$$\pi^*(s) = \operatorname{argmax}_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

- For every state s and every policy π

$$V^*(s) = V^{\pi^*}(s) \geq V^{\pi}(s)$$

Value Iteration & Policy Iteration

- Note that the value function and policy are correlated

$$V^{\pi}(s) = R(s) + \gamma \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^{\pi}(s')$$

$$\pi(s) = \operatorname{argmax}_{a \in A} \sum_{s' \in S} P_{sa}(s') V^{\pi}(s')$$

- It is feasible to perform iterative update towards the optimal value function and optimal policy
 - Value iteration
 - Policy iteration

Value Iteration

- For an MDP with finite state and action spaces
 $|S| < \infty, |A| < \infty$

- Value iteration is performed as

1. For each state s , initialize $V(s) = 0$

2. Repeat until convergence {

For each state, update

$$V(s) = R(s) + \gamma \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V(s')$$

}

- Note that there is no explicit policy in above calculation

Value Iteration Example: Shortest Path

g			

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

V_1

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

V_2

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

V_3

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

V_4

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

V_5

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

V_6

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

V_7

Policy Iteration

- For an MDP with finite state and action spaces
 $|S| < \infty, |A| < \infty$

- Policy iteration is performed as



1. Initialize π randomly
2. Repeat until convergence {
 - a) Let $V := V^\pi$
 - b) For each state, update

$$\pi(s) = \operatorname{argmax}_{a \in A} \sum_{s' \in S} P_{sa}(s') V^\pi(s')$$


}

- The step of value function update could be time-consuming

Policy Iteration



- Policy evaluation
 - Estimate V^π
 - Iterative policy evaluation
- Policy improvement
 - Generate $\pi' \geq \pi$
 - Greedy policy improvement



Value Iteration vs. Policy Iteration

- Value iteration

1. For each state s , initialize $V(s) = 0$
2. Repeat until convergence {
For each state, update

$$V(s) = R(s) + \gamma \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V(s')$$

}

- Policy iteration

1. Initialize π randomly
2. Repeat until convergence {
 - a) Let $V := V^\pi$
 - b) For each state, update

$$\pi(s) = \operatorname{argmax}_{a \in A} \sum_{s' \in S} P_{sa}(s') V^\pi(s')$$

}

Remarks:

1. Value iteration is a greedy update strategy
2. In policy iteration, the value function update by bellman equation is costly
3. For small-space MDPs, policy iteration is often very fast and converges quickly
4. For large-space MDPs, value iteration is more practical (efficient)
5. If there is no state-transition loop, it is better to use value iteration

Learning an MDP Model

- So far we have been focused on
 - Calculating the optimal value function
 - Learning the optimal policygiven a known MDP model
 - i.e. the state transition $P_{sa}(s')$ and reward function $R(s)$ are explicitly given
- In realistic problems, often the state transition and reward function are not explicitly given
- We only have some observations by experience
 - e.g., play games, inventory management with unknown demand, online advertisement...

Learning an MDP Model

- Learn an MDP model from “experience”

- Learning state transition probabilities $P_{sa}(s')$

$$P_{sa}(s') = \frac{\text{\#times we took action } a \text{ in state } s \text{ and got to state } s'}{\text{\#times we took action } a \text{ in state } s}$$

- Learning reward $R(s)$, i.e. the expected immediate reward

$$R(s) = \text{avg}(R(s)^{(i)})$$

Learning Model and Optimizing Policy

Algorithm

1. Initialize π randomly
2. Repeat until convergence {
 - a) Execute π in the MDP for some number of trials
 - b) Using the accumulated experience in the MDP, update our estimates for P_{sa} and R
 - c) Apply value iteration with the estimated P_{sa} and R to get the new estimated value function V
 - d) Update π to be the greedy policy w.r.t. V}

- Another branch of solution is to directly learning value & policy from experience without building an MDP
- i.e. **Model-free Reinforcement Learning**

Model-free RL

Model-free Prediction

Model-free Reinforcement Learning

- In realistic problems, often the state transition and reward function are not explicitly given
- Model-free RL is to directly learn value & policy from experience without building an MDP
- Key steps: (1) estimate value function; (2) optimize policy

Value Function Estimation




- In model-based RL (MDP), the value function is calculated by dynamic programming

$$\begin{aligned} V^\pi(s) &= \mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots | s_0 = s, \pi] \\ &= R(s) + \gamma \sum_{s' \in \mathcal{S}} P_{s\pi(s)}(s') V^\pi(s') \end{aligned}$$

- Now in model-free RL
 - We cannot directly know P_{sa} and R
 - But we have a list of experiences to estimate the values

Monte-Carlo Methods



- Monte-Carlo methods are a broad class of computational algorithms that rely on repeated random sampling to obtain numerical results.
- Example, to calculate the circle's surface



$$\text{Circle Surface} = \text{Square Surface} \times \frac{\text{\#points in circle}}{\text{\#points in total}}$$

Monte-Carlo Methods

- Go: to estimate the winning rate given the current state



$$\text{Win Rate } (s) = \frac{\text{\#win simulation cases started from } s}{\text{\#simulation cases started from } s \text{ in total}}$$

Monte-Carlo Value Estimation

- Goal: learn V^π from experience under policy π
- Recall that the return is the total discounted reward

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots \gamma^{T-1} R_T$$

- Recall that the value function is the expected return

$$V^\pi(s) = \mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots | s_0 = s, \pi]$$

$$= \mathbb{E}[G_t | s_t = s, \pi]$$

$$\simeq \frac{1}{N} \sum_{i=1}^N G_t^{(i)}$$

- Sample N episodes from state s using policy π
 - Calculate the average of cumulative reward
 - $G_t^{(i)}$ is the G_t of i th sample
- Monte-Carlo policy evaluation uses empirical mean return instead of expected return

Monte-Carlo Value Estimation

- Implementation
 - Sample episodes under policy π

$$s_0^{(i)} \xrightarrow[R_1^{(i)}]{a_0^{(i)}} s_1^{(i)} \xrightarrow[R_2^{(i)}]{a_1^{(i)}} s_2^{(i)} \xrightarrow[R_3^{(i)}]{a_2^{(i)}} s_3^{(i)} \dots s_T^{(i)} \sim \pi$$

- Every time-step t that state s is visited in an episode
 - Increment counter $N(s) = N(s) + 1$
 - Increment total return $S(s) = S(s) + G_t$
 - Value is estimated by mean return $V(s) = S(s)/N(s)$
 - By law of large numbers
$$V(s) \rightarrow V^{\pi}(s) \text{ as } N(s) \rightarrow \infty$$

Incremental Monte-Carlo Updates

- Update $V(s)$ incrementally after each episode
- For each state S_t with cumulative return G_t

$$N(s_t) = N(s_t) + 1$$

$$V(s_t) = V(s_t) + \frac{1}{N(s_t)} (G_t - V(s_t))$$

- For non-stationary problems (i.e. the environment could be varying over time), it can be useful to track a running mean, i.e. forget old episodes

$$V(s_t) = V(s_t) + \alpha(G_t - V(s_t))$$

Monte-Carlo Value Estimation

$$\text{Idea: } V(s) = \frac{1}{N} \sum_{i=1}^N G_t^{(i)}$$

$$\text{Implementation: } V(s_t) = V(s_t) + \alpha(G_t - V(s_t))$$

- MC methods learn directly from episodes of experience
- MC is model-free: no knowledge of MDP transitions / rewards
- MC learns from **complete** episodes
- MC uses the simplest possible idea: value = mean return
- Caveat: can only apply MC to episodic MDPs
 - All episodes must terminate

Lecture 15 Wrap-up

- ✓ Reinforcement Learning
- ✓ Model-based Reinforcement Learning
- ✓ Model-free Reinforcement Learning

Next Lecture

- Supervised learning
 - Linear regression
 - Logistic regression
 - SVM and kernel
 - Tree models
- Deep learning
 - Neural networks
 - Convolutional NN
 - Recurrent NN
- Unsupervised learning
 - Clustering
 - PCA (Dimension Reduction)
 - EM
- Reinforcement learning
 - MDP
 - ADP
 - Deep Q-Network



Questions?

Shan Wang (王杉)

<https://wangshan731.github.io/>